



RKeOps



KERNEL OPERATIONS WITH SYMBOLIC TENSORS ON THE GPU IN R

Ghislain Durif² & Benjamin Charlier¹ & Chloé Serre-Combe¹ & Amélie Vernay^{1,*}

June 21st 2023 – Rencontres R 2023

¹IMAG, Université de Montpellier, CNRS UMR 5149, Montpellier, France

²LBMC, ENS de Lyon, CNRS UMR 5239, Inserm U1293, Université Claude Bernard Lyon 1, Lyon, France

*amelie.vernay@umontpellier.fr

A motivating example

For $i = 1, \dots, M$ we want to compute the **reduction**

$$a_i = \sum_{j=1}^N K(\mathbf{x}_i, \mathbf{y}_j) b_j, \quad (1)$$

with

- source points $\mathbf{y}_1, \dots, \mathbf{y}_N \in \mathbb{R}^D$ with associated weights $b_1, \dots, b_N \in \mathbb{R}$
- target points $\mathbf{x}_1, \dots, \mathbf{x}_M \in \mathbb{R}^D$
- a Gaussian kernel $K(\mathbf{x}_i, \mathbf{y}_j) = \exp(-\|\mathbf{x}_i - \mathbf{y}_j\|_2^2)$

A motivating example

For $i = 1, \dots, M$ we want to compute the **reduction**

$$a_i = \sum_{j=1}^N K(\mathbf{x}_i, \mathbf{y}_j) b_j, \quad (1)$$

with

- source points $\mathbf{y}_1, \dots, \mathbf{y}_N \in \mathbb{R}^D$ with associated weights $b_1, \dots, b_N \in \mathbb{R}$
- target points $\mathbf{x}_1, \dots, \mathbf{x}_M \in \mathbb{R}^D$
- a Gaussian kernel $K(\mathbf{x}_i, \mathbf{y}_j) = \exp(-\|\mathbf{x}_i - \mathbf{y}_j\|_2^2)$

Limitations of basic routines

- computation of all the elements $(K(\mathbf{x}_i, \mathbf{y}_j))_{i,j} \rightarrow \mathcal{O}(MN)$ time complexity
- storage as a dense $M \times N$ matrix $\rightarrow \mathcal{O}(MN)$ memory usage

A motivating example

For $i = 1, \dots, M$ we want to compute the **reduction**

$$a_i = \sum_{j=1}^N K(\mathbf{x}_i, \mathbf{y}_j) b_j, \quad (1)$$

with

- source points $\mathbf{y}_1, \dots, \mathbf{y}_N \in \mathbb{R}^D$ with associated weights $b_1, \dots, b_N \in \mathbb{R}$
- target points $\mathbf{x}_1, \dots, \mathbf{x}_M \in \mathbb{R}^D$
- a Gaussian kernel $K(\mathbf{x}_i, \mathbf{y}_j) = \exp(-\|\mathbf{x}_i - \mathbf{y}_j\|_2^2)$

Limitations of basic routines

- computation of all the elements $(K(\mathbf{x}_i, \mathbf{y}_j))_{i,j} \rightarrow \mathcal{O}(MN)$ time complexity
- storage as a dense $M \times N$ matrix $\rightarrow \mathcal{O}(MN)$ memory usage

\rightarrow Impossible in high dimension! \leftarrow

KeOps: Kernel Operations

RKeOps: Kernel Operations in R

RKeOps: Kernel Operations in R

- ▶ Perform **fast reductions** of very large arrays ($M, N \simeq 10^6$),

RKeOps: Kernel Operations in R

- ▶ Perform **fast reductions** of very large arrays ($M, N \simeq 10^6$),
- ▶ with **effortless computation on GPU** without memory overflow,

RKeOps: Kernel Operations in R

- ▶ Perform **fast reductions** of very large arrays ($M, N \simeq 10^6$),
- ▶ with **effortless computation on GPU** without memory overflow,
- ▶ with **automatic differentiation** up to arbitrary orders.


RKeOps: Kernel Operations in R

- ▶ Perform **fast reductions** of very large arrays ($M, N \simeq 10^6$),
- ▶ with **effortless computation on GPU** without memory overflow,
- ▶ with **automatic differentiation** up to arbitrary orders.


In short, RKeOps offers gains on runtime and memory usage by performing on-the-fly compilation with **symbolic matrices**.

→ <https://www.kernel-operations.io> ←


Symbolic evaluation with RKeOps LazyTensors



Dense matrix



Sparse matrix




Symbolic matrix


Symbolic evaluation with RKeOps LazyTensors



Dense matrix



Sparse matrix




Symbolic matrix

Symbolic evaluation with RKeOps LazyTensors



Symbolic matrices: Matrices whose entries are given by a mathematical formula.


Symbolic evaluation with RKeOps LazyTensors



Symbolic matrices: Matrices whose entries are given by a mathematical formula.

RKeOps LazyTensors: Wrappers around R data arrays that embody symbolic matrices.

Symbolic evaluation with RKeOps LazyTensors



Symbolic matrices: Matrices whose entries are given by a mathematical formula.

RKeOps LazyTensors: Wrappers around R data arrays that embody symbolic matrices.

1. import RKeOps
2. create **LazyTensors** with your data
3. perform *any* kinds of reduction using friendly R native syntax

Symbolic evaluation with RKeOps LazyTensors

Let us perform the Gaussian reduction (1) with RKeOps: $\left[\sum_{j=1}^N k(\mathbf{x}_i, \mathbf{y}_j) b_j \right]_{i=1}^M$

```
# Create large point clouds
N <- 10^5; D <- 15
x <- matrix(rnorm(N*D), N, D)
y <- matrix(rnorm(N*D), N, D)
b <- matrix(rnorm(N*D), N, D)

# Turn dense arrays into symbolic matrices
x_i <- LazyTensor(x, "i")
y_j <- LazyTensor(y, "j")
b_j <- LazyTensor(b, "j")
K_ij <- exp( -sum((x_i - y_j)^2) )    # symbolic N*N Gaussian kernel

# Call sum() reduction to trigger the computation
a_i <- sum(K_ij * b_j, index = "j")
dim(a_i); class(a_i)
# [1] 100000      15
# [1] "matrix" "array"
```

Total running time of the script: 4.373 sec.¹

1. using 16 cores of an Intel Xeon Gold 6142 processor

Symbolic evaluation with RKeOps LazyTensors

Let us perform the Gaussian reduction (1) with RKeOps: $\left[\sum_{j=1}^N k(\mathbf{x}_i, \mathbf{y}_j) b_j \right]_{i=1}^M$

```
# Create large point clouds
N <- 10^5; D <- 15
x <- matrix(rnorm(N*D), N, D)
y <- matrix(rnorm(N*D), N, D)
b <- matrix(rnorm(N*D), N, D)

# Turn dense arrays into symbolic matrices
x_i <- LazyTensor(x, "i")
y_j <- LazyTensor(y, "j")
b_j <- LazyTensor(b, "j")
K_ij <- exp( -sum((x_i - y_j)^2) )    # symbolic N*N Gaussian kernel

# Call sum() reduction to trigger the computation
a_i <- sum(K_ij * b_j, index = "j")
dim(a_i); class(a_i)
# [1] 100000      15
# [1] "matrix" "array"
```

Total running time of the script: 4.373 sec.¹

1. using 16 cores of an Intel Xeon Gold 6142 processor

Symbolic evaluation with RKeOps LazyTensors

Let us perform the Gaussian reduction (1) with RKeOps: $\left[\sum_{j=1}^N k(\mathbf{x}_i, \mathbf{y}_j) b_j \right]_{i=1}^M$

```
# Create large point clouds
N <- 10^5; D <- 15
x <- matrix(rnorm(N*D), N, D)
y <- matrix(rnorm(N*D), N, D)
b <- matrix(rnorm(N*D), N, D)

# Turn dense arrays into symbolic matrices
x_i <- LazyTensor(x, "i")
y_j <- LazyTensor(y, "j")
b_j <- LazyTensor(b, "j")
K_ij <- exp( -sum((x_i - y_j)^2) ) # symbolic N×N Gaussian kernel

# Call sum() reduction to trigger the computation
a_i <- sum(K_ij * b_j, index = "j")
dim(a_i); class(a_i)
# [1] 100000      15
# [1] "matrix" "array"
```

Total running time of the script: 4.373 sec.¹

1. using 16 cores of an Intel Xeon Gold 6142 processor

Symbolic evaluation with RKeOps LazyTensors

Let us perform the Gaussian reduction (1) with RKeOps: $\left[\sum_{j=1}^N k(\mathbf{x}_i, \mathbf{y}_j) b_j \right]_{i=1}^M$

```
# Create large point clouds
N <- 10^5; D <- 15
x <- matrix(rnorm(N*D), N, D)
y <- matrix(rnorm(N*D), N, D)
b <- matrix(rnorm(N*D), N, D)

# Turn dense arrays into symbolic matrices
x_i <- LazyTensor(x, "i")
y_j <- LazyTensor(y, "j")
b_j <- LazyTensor(b, "j")
K_ij <- exp( -sum((x_i - y_j)^2) ) # symbolic N×N Gaussian kernel

# Call sum() reduction to trigger the computation
a_i <- sum(K_ij * b_j, index = "j")
dim(a_i); class(a_i)
# [1] 100000      15
# [1] "matrix" "array"
```

Total running time of the script: 4.373 sec.¹

1. using 16 cores of an Intel Xeon Gold 6142 processor

Generic reduction with RKeOps

RKeOps supports all kinds of reduction. For $1 \leq i \leq M$, compute

$$\left[\text{Reduction}_{j=1,\dots,M} F(\mathbf{p}^1, \mathbf{p}^2, \dots, \mathbf{x}_i^1, \mathbf{x}_i^2, \dots, \mathbf{y}_j^1, \mathbf{y}_j^2, \dots) \right]_{i=1,\dots,M}$$

where

- ▶ "Reduction" can be any reduction over a dimension (Sum, Max, ArgMax, LogSumExp...)
- ▶ F is a vector-valued formula
- ▶ $\mathbf{x}_i^1, \mathbf{x}_i^2, \dots$ are vector variables indexed by i
- ▶ $\mathbf{y}_j^1, \mathbf{y}_j^2, \dots$ are vector variables indexed by j
- ▶ $\mathbf{p}^1, \mathbf{p}^2, \dots$ are vector parameter fixed across indices

2. https://github.com/getkeops/keops/blob/main/rkeops/vignettes/LazyTensor_rkeops.Rmd

Generic reduction with RKeOps

RKeOps supports all kinds of reduction. For $1 \leq i \leq M$, compute

$$\left[\text{Reduction}_{j=1,\dots,M} F(\mathbf{p}^1, \mathbf{p}^2, \dots, \mathbf{x}_i^1, \mathbf{x}_i^2, \dots, \mathbf{y}_j^1, \mathbf{y}_j^2, \dots) \right]_{i=1,\dots,M}$$

where

- ▶ "Reduction" can be any reduction over a dimension (Sum, Max, ArgMax, LogSumExp...)
- ▶ F is a vector-valued formula
- ▶ $\mathbf{x}_i^1, \mathbf{x}_i^2, \dots$ are vector variables indexed by i
- ▶ $\mathbf{y}_j^1, \mathbf{y}_j^2, \dots$ are vector variables indexed by j
- ▶ $\mathbf{p}^1, \mathbf{p}^2, \dots$ are vector parameter fixed across indices

The full range of reductions and operations provided by RKeOps is available in the vignette².

2. https://github.com/getkeops/keops/blob/main/rkeops/vignettes/LazyTensor_rkeops.Rmd

Generic reduction with RKeOps

RKeOps supports all kinds of reduction. For $1 \leq i \leq M$, compute

$$\left[\text{Reduction}_{j=1,\dots,M} F(\mathbf{p}^1, \mathbf{p}^2, \dots, \mathbf{x}_i^1, \mathbf{x}_i^2, \dots, \mathbf{y}_j^1, \mathbf{y}_j^2, \dots) \right]_{i=1,\dots,M}$$

where

- ▶ "Reduction" can be any reduction over a dimension (Sum, Max, ArgMax, LogSumExp...)
- ▶ F is a vector-valued formula
- ▶ $\mathbf{x}_i^1, \mathbf{x}_i^2, \dots$ are vector variables indexed by i
- ▶ $\mathbf{y}_j^1, \mathbf{y}_j^2, \dots$ are vector variables indexed by j
- ▶ $\mathbf{p}^1, \mathbf{p}^2, \dots$ are vector parameter fixed across indices


The full range of reductions and operations provided by RKeOps is available in the vignette².

Note: RKeOps also supports reductions on complex data!

2. https://github.com/getkeops/keops/blob/main/rkeops/vignettes/LazyTensor_rkeops.Rmd

Example I - K-means clustering

At each iteration, compute $\operatorname{argmin}_{j=1,\dots,K} \|x_i - c_j\|$ for $i = 1, \dots, N$, where c_j is the centroid of cluster j .



Example of 50-means clustering with $N = 10^5$ points in \mathbb{R}^2 and the Euclidean distance.


Time for 10 iterations on CPU: 2.224 sec. (0.222 sec. per iteration)

Example II - Kernel interpolation

For $\lambda \in \mathbb{R}_+$, let us solve a linear system of the form

$$\mathbf{a}^* = \underset{\mathbf{a}}{\operatorname{argmin}} \frac{1}{2} \langle \mathbf{a}, (\lambda \operatorname{Id} + \mathbf{K}) \mathbf{a} \rangle - \langle \mathbf{a}, \mathbf{b} \rangle = (\lambda \operatorname{Id} + \mathbf{K})^{-1} \mathbf{b}$$

where \mathbf{K} is a symmetric, positive definite linear operator defined with a symbolic formula.



Example of 1D interpolation with a Gaussian kernel matrix on $N = 10^4$ points.

Time to perform the interpolation with a precision of 10^{-6} on CPU: 13 sec.

AutoDiff with RKeOps

What if we need the gradient of $\mathbf{a} = (a_1, \dots, a_M)$, say with respect to \mathbf{y} ?

$$\partial_{\mathbf{y}} \mathbf{a} = \partial_{\mathbf{y}} (a_i)_{i=1, \dots, M} = \partial_{\mathbf{y}} \left(\sum_{j=1}^N k(\mathbf{x}_i, \mathbf{y}_j) b_j \right)_{i=1, \dots, M}$$

RKeOps LazyTensors → do not support automatic differentiation (yet!)

KeOps → provides an autodiff engine for formulae wrapped in `Grad()`:

```
# Define a formula with a gradient
formula_grad <- 'Grad(Sum_Reduction(Exp(-SqDist(x,y)) * b , 0), y, e)'
variables <- c('x = Vi(15)', 'y = Vj(15)', 'b = Vj(15)', 'e = Vi(15)')
# Compile the corresponding operator
gaussian_kernel_grad <- keops_kernel(formula_grad, variables)
# Declare a new tensor used as the input of the gradient operator
e <- matrix(rnorm(N*D), N, D)
# Computation
res <- gaussian_kernel_grad(list(x, y, b, e))
```

Total running time of the script: 11.975 sec.³

3. using 16 cores of an Intel Xeon Gold 6142 processor

AutoDiff with RKeOps

What if we need the gradient of $\mathbf{a} = (a_1, \dots, a_M)$, say with respect to \mathbf{y} ?

$$\partial_{\mathbf{y}} \mathbf{a} = \partial_{\mathbf{y}} (a_i)_{i=1, \dots, M} = \partial_{\mathbf{y}} \left(\sum_{j=1}^N k(\mathbf{x}_i, \mathbf{y}_j) b_j \right)_{i=1, \dots, M}$$

RKeOps LazyTensors → do not support automatic differentiation (**yet!**)

KeOps → provides an autodiff engine for formulae wrapped in `Grad()`:

```
# Define a formula with a gradient
formula_grad <- 'Grad(Sum_Reduction(Exp(-SqDist(x,y)) * b , 0), y, e)'
variables <- c('x = Vi(15)', 'y = Vj(15)', 'b = Vj(15)', 'e = Vi(15)')
# Compile the corresponding operator
gaussian_kernel_grad <- keops_kernel(formula_grad, variables)
# Declare a new tensor used as the input of the gradient operator
e <- matrix(rnorm(N*D), N, D)
# Computation
res <- gaussian_kernel_grad(list(x, y, b, e))
```

Total running time of the script: 11.975 sec.³

3. using 16 cores of an Intel Xeon Gold 6142 processor

AutoDiff with RKeOps

What if we need the gradient of $\mathbf{a} = (a_1, \dots, a_M)$, say with respect to \mathbf{y} ?

$$\partial_{\mathbf{y}} \mathbf{a} = \partial_{\mathbf{y}} (a_i)_{i=1, \dots, M} = \partial_{\mathbf{y}} \left(\sum_{j=1}^N k(\mathbf{x}_i, \mathbf{y}_j) b_j \right)_{i=1, \dots, M}$$

RKeOps LazyTensors → do not support automatic differentiation (**yet!**)

KeOps → provides an autodiff engine for formulae wrapped in `Grad()`:

```
# Define a formula with a gradient
formula_grad <- 'Grad(Sum_Reduction(Exp(-SqDist(x,y)) * b , 0), y, e)'
variables <- c('x = Vi(15)', 'y = Vj(15)', 'b = Vj(15)', 'e = Vi(15)')
# Compile the corresponding operator
gaussian_kernel_grad <- keops_kernel(formula_grad, variables)
# Declare a new tensor used as the input of the gradient operator
e <- matrix(rnorm(N*D), N, D)
# Computation
res <- gaussian_kernel_grad(list(x, y, b, e))
```

Total running time of the script: 11.975 sec.³

3. using 16 cores of an Intel Xeon Gold 6142 processor

AutoDiff with RKeOps

What if we need the gradient of $\mathbf{a} = (a_1, \dots, a_M)$, say with respect to \mathbf{y} ?

$$\partial_{\mathbf{y}} \mathbf{a} = \partial_{\mathbf{y}} (a_i)_{i=1, \dots, M} = \partial_{\mathbf{y}} \left(\sum_{j=1}^N k(\mathbf{x}_i, \mathbf{y}_j) b_j \right)_{i=1, \dots, M}$$

RKeOps LazyTensors → do not support automatic differentiation (yet!)

KeOps → provides an autodiff engine for formulae wrapped in `Grad()`:

```
# Define a formula with a gradient
formula_grad <- 'Grad(Sum_Reduction(Exp(-SqDist(x,y)) * b , 0), y, e)'
variables <- c('x = Vi(15)', 'y = Vj(15)', 'b = Vj(15)', 'e = Vi(15)')
# Compile the corresponding operator
gaussian_kernel_grad <- keops_kernel(formula_grad, variables)
# Declare a new tensor used as the input of the gradient operator
e <- matrix(rnorm(N*D), N, D)
# Computation
res <- gaussian_kernel_grad(list(x, y, b, e))
```

Total running time of the script: 11.975 sec.³

3. using 16 cores of an Intel Xeon Gold 6142 processor

AutoDiff with RKeOps

What if we need the gradient of $\mathbf{a} = (a_1, \dots, a_M)$, say with respect to \mathbf{y} ?

$$\partial_{\mathbf{y}} \mathbf{a} = \partial_{\mathbf{y}} (a_i)_{i=1, \dots, M} = \partial_{\mathbf{y}} \left(\sum_{j=1}^N k(\mathbf{x}_i, \mathbf{y}_j) b_j \right)_{i=1, \dots, M}$$

RKeOps LazyTensors → do not support automatic differentiation (yet!)

KeOps → provides an autodiff engine for formulae wrapped in `Grad()`:

```
# Define a formula with a gradient
formula_grad <- 'Grad(Sum_Reduction(Exp(-SqDist(x,y)) * b , 0), y, e)'
variables <- c('x = Vi(15)', 'y = Vj(15)', 'b = Vj(15)', 'e = Vi(15)')
# Compile the corresponding operator
gaussian_kernel_grad <- keops_kernel(formula_grad, variables)
# Declare a new tensor used as the input of the gradient operator
e <- matrix(rnorm(N*D), N, D)
# Computation
res <- gaussian_kernel_grad(list(x, y, b, e))
```

Total running time of the script: 11.975 sec.³

3. using 16 cores of an Intel Xeon Gold 6142 processor

AutoDiff with RKeOps

What if we need the gradient of $\mathbf{a} = (a_1, \dots, a_M)$, say with respect to \mathbf{y} ?

$$\partial_{\mathbf{y}} \mathbf{a} = \partial_{\mathbf{y}} (a_i)_{i=1, \dots, M} = \partial_{\mathbf{y}} \left(\sum_{j=1}^N k(\mathbf{x}_i, \mathbf{y}_j) b_j \right)_{i=1, \dots, M}$$

RKeOps LazyTensors → do not support automatic differentiation (yet!)

KeOps → provides an autodiff engine for formulae wrapped in `Grad()`:

```
# Define a formula with a gradient
formula_grad <- 'Grad(Sum_Reduction(Exp(-SqDist(x,y)) * b , 0), y, e)'
variables <- c('x = Vi(15)', 'y = Vj(15)', 'b = Vj(15)', 'e = Vi(15)')
# Compile the corresponding operator
gaussian_kernel_grad <- keops_kernel(formula_grad, variables)
# Declare a new tensor used as the input of the gradient operator
e <- matrix(rnorm(N*D), N, D)
# Computation
res <- gaussian_kernel_grad(list(x, y, b, e))
```

Total running time of the script: 11.975 sec.³

3. using 16 cores of an Intel Xeon Gold 6142 processor

AutoDiff with RKeOps

What if we need the gradient of $\mathbf{a} = (a_1, \dots, a_M)$, say with respect to \mathbf{y} ?

$$\partial_{\mathbf{y}} \mathbf{a} = \partial_{\mathbf{y}} (a_i)_{i=1, \dots, M} = \partial_{\mathbf{y}} \left(\sum_{j=1}^N k(\mathbf{x}_i, \mathbf{y}_j) b_j \right)_{i=1, \dots, M}$$

RKeOps LazyTensors → do not support automatic differentiation (yet!)

KeOps → provides an autodiff engine for formulae wrapped in `Grad()`:

```
# Define a formula with a gradient
formula_grad <- 'Grad(Sum_Reduction(Exp(-SqDist(x,y)) * b , 0), y, e)'
variables <- c('x = Vi(15)', 'y = Vj(15)', 'b = Vj(15)', 'e = Vi(15)')

# Compile the corresponding operator
gaussian_kernel_grad <- keops_kernel(formula_grad, variables)

# Declare a new tensor used as the input of the gradient operator
e <- matrix(rnorm(N*D), N, D)

# Computation
res <- gaussian_kernel_grad(list(x, y, b, e))
```

Total running time of the script: 11.975 sec.³

3. using 16 cores of an Intel Xeon Gold 6142 processor

AutoDiff with RKeOps

What if we need the gradient of $\mathbf{a} = (a_1, \dots, a_M)$, say with respect to \mathbf{y} ?

$$\partial_{\mathbf{y}} \mathbf{a} = \partial_{\mathbf{y}} (a_i)_{i=1, \dots, M} = \partial_{\mathbf{y}} \left(\sum_{j=1}^N k(\mathbf{x}_i, \mathbf{y}_j) b_j \right)_{i=1, \dots, M}$$

RKeOps LazyTensors → do not support automatic differentiation (yet!)

KeOps → provides an autodiff engine for formulae wrapped in `Grad()`:

```
# Define a formula with a gradient
formula_grad <- 'Grad(Sum_Reduction(Exp(-SqDist(x,y)) * b , 0), y, e)'
variables <- c('x = Vi(15)', 'y = Vj(15)', 'b = Vj(15)', 'e = Vi(15)')
# Compile the corresponding operator
gaussian_kernel_grad <- keops_kernel(formula_grad, variables)
# Declare a new tensor used as the input of the gradient operator
e <- matrix(rnorm(N*D), N, D)
# Computation
res <- gaussian_kernel_grad(list(x, y, b, e))
```

Total running time of the script: 11.975 sec.³

3. using 16 cores of an Intel Xeon Gold 6142 processor

AutoDiff with RKeOps

What if we need the gradient of $\mathbf{a} = (a_1, \dots, a_M)$, say with respect to \mathbf{y} ?

$$\partial_{\mathbf{y}} \mathbf{a} = \partial_{\mathbf{y}} (a_i)_{i=1, \dots, M} = \partial_{\mathbf{y}} \left(\sum_{j=1}^N k(\mathbf{x}_i, \mathbf{y}_j) b_j \right)_{i=1, \dots, M}$$

RKeOps LazyTensors → do not support automatic differentiation (yet!)

KeOps → provides an autodiff engine for formulae wrapped in `Grad()`:

```
# Define a formula with a gradient
formula_grad <- 'Grad(Sum_Reduction(Exp(-SqDist(x,y)) * b , 0), y, e)'
variables <- c('x = Vi(15)', 'y = Vj(15)', 'b = Vj(15)', 'e = Vi(15)')
# Compile the corresponding operator
gaussian_kernel_grad <- keops_kernel(formula_grad, variables)
# Declare a new tensor used as the input of the gradient operator
e <- matrix(rnorm(N*D), N, D)
# Computation
res <- gaussian_kernel_grad(list(x, y, b, e))
```

Total running time of the script: 11.975 sec.³

3. using 16 cores of an Intel Xeon Gold 6142 processor

AutoDiff with RKeOps

What if we need the gradient of $\mathbf{a} = (a_1, \dots, a_M)$, say with respect to \mathbf{y} ?

$$\partial_{\mathbf{y}} \mathbf{a} = \partial_{\mathbf{y}} (a_i)_{i=1, \dots, M} = \partial_{\mathbf{y}} \left(\sum_{j=1}^N k(\mathbf{x}_i, \mathbf{y}_j) b_j \right)_{i=1, \dots, M}$$

RKeOps LazyTensors → do not support automatic differentiation (**yet!**)

KeOps → provides an autodiff engine for formulae wrapped in `Grad()`:

```
# Define a formula with a gradient
formula_grad <- 'Grad(Sum_Reduction(Exp(-SqDist(x,y)) * b , 0), y, e)'
variables <- c('x = Vi(15)', 'y = Vj(15)', 'b = Vj(15)', 'e = Vi(15)')
# Compile the corresponding operator
gaussian_kernel_grad <- keops_kernel(formula_grad, variables)
# Declare a new tensor used as the input of the gradient operator
e <- matrix(rnorm(N*D), N, D)
# Computation
res <- gaussian_kernel_grad(list(x, y, b, e))
```

Total running time of the script: 11.975 sec.³

3. using 16 cores of an Intel Xeon Gold 6142 processor

KeOps⁴ Core library

- written in C++/Python
- dependencies: a C++ compiler (`g++`, `clang`) and `nvrtc` headers provided by `CUDA` for GPU computing

4. Benjamin Charlier et al. “Kernel Operations on the GPU, with Autodiff, without Memory Overflows”. In: *Journal of Machine Learning Research* 22.74 (2021), pp. 1–6. URL: <http://jmlr.org/papers/v22/20-275.html>.

5. Kevin Ushey, JJ Allaire, and Yuan Tang. *reticulate: Interface to 'Python'*. <https://rstudio.github.io/reticulate/>, <https://github.com/rstudio/reticulate>. 2023.

KeOps⁴ Core library

- written in C++/Python
- dependencies: a C++ compiler (`g++`, `clang`) and `nvrtc` headers provided by `CUDA` for GPU computing

PyKeOps Python binders for KeOps (both NumPy and PyTorch)

4. Benjamin Charlier et al. “Kernel Operations on the GPU, with Autodiff, without Memory Overflows”. In: *Journal of Machine Learning Research* 22.74 (2021), pp. 1–6. URL: <http://jmlr.org/papers/v22/20-275.html>.

5. Kevin Ushey, JJ Allaire, and Yuan Tang. *reticulate: Interface to 'Python'*. <https://rstudio.github.io/reticulate/>, <https://github.com/rstudio/reticulate>. 2023.

KeOps⁴ Core library

- written in C++/Python
- dependencies: a C++ compiler (g++, clang) and nvrtc headers provided by CUDA for GPU computing

PyKeOps Python binders for KeOps (both NumPy and PyTorch)

RKeOps R binder for KeOps

- since v.2.0, directly uses PyKeOps through reticulate⁵
- GPU computing directly inside R: just type `rkeops_use_gpu()`!
- soon on the CRAN, already available on github:


```
install.packages("remotes")  
remotes::install_github("getkeops/keops", subdir = "rkeops")
```

4. Benjamin Charlier et al. "Kernel Operations on the GPU, with Autodiff, without Memory Overflows". In: *Journal of Machine Learning Research* 22.74 (2021), pp. 1–6. URL: <http://jmlr.org/papers/v22/20-275.html>.

5. Kevin Ushey, JJ Allaire, and Yuan Tang. *reticulate: Interface to 'Python'*. <https://rstudio.github.io/reticulate/>, <https://github.com/rstudio/reticulate>. 2023.

Benchmark I - Gaussian convolution


Runtimes⁶ for Gaussian convolution with N samples in \mathbb{R}^{15} .



6. CPU: 16 cores of an Intel Xeon Gold 6142 processor. GPU: Nvidia A10.

Benchmark II - KNN-search

Runtimes⁷ for 10-Nearest Neighbors search with N samples in \mathbb{R}^3 .



7. CPU: 16 cores of an Intel Xeon Gold 6142 processor. GPU: Nvidia A10.

RKeOps: Fast kernel operations on GPU without memory overflow and with automatic differentiation, directly **inside R**

- ▶ full documentation, tutorials, examples, benchmarks and more at <https://www.kernel-operations.io>
- ▶ active development and open contributions at <https://github.com/getkeops/keops/blob/main/rkeops/>
- ▶ current install process:

```
install.packages("remotes")
remotes::install_github("getkeops/keops", subdir = "rkeops")
```
- ▶ available on CRAN soon!

Thank you for your attention! Questions?

- [1] Benjamin Charlier et al. “Kernel Operations on the GPU, with Autodiff, without Memory Overflows”. In: *Journal of Machine Learning Research* 22.74 (2021), pp. 1–6.
URL: <http://jmlr.org/papers/v22/20-275.html>.
- [2] Kevin Ushey, JJ Allaire, and Yuan Tang. *reticulate: Interface to 'Python'*.
<https://rstudio.github.io/reticulate/>, <https://github.com/rstudio/reticulate>. 2023.